

PRODUCT REQUIREMENTS DOCUMENT • v2.0 • COMPLETE

TrustGraph

The Trust Layer of the Internet

5 Phases • 36-Month Build Timeline • Full Feature Architecture

Proof over Claims

Transparent Scoring

Portable Reputation

Evidence-First

Privacy-Safe

1. Vision & Core Thesis

"The credit score of human capability — except transparent and explainable."

TrustGraph is the missing layer of the internet. Every major coordination problem the internet has solved — search, communication, payments, code — had a winner. Reputation portability has no winner yet. TrustGraph is built to be that winner.

The product aggregates digital evidence across platforms, transforms it into a transparent, explainable trust profile, and makes that profile portable — one link that answers the question every professional relationship starts with:

"Can I trust this person?"

What the Internet Built — and the Gap It Left

Solved

- Information → Search engines
- Communication → Social media
- Payments → Stripe / PayPal
- Code collaboration → GitHub
- Identity → Auth0 / Clerk

Never Solved

Reputation Portability

Your reputation is fragmented across every platform you've ever worked on. It cannot move with you. It resets with every new context.

2. The Problem Space

Reputation is fragmented. Every platform is an island. TrustGraph bridges them.

For Recruiters

Cannot distinguish real builders from resume optimizers. Keyword-matching and claim-checking is the state of the art. No signal for actual output.

For Founders

No reliable way to evaluate potential co-founders, contributors, or early employees beyond gut feel and LinkedIn endorsements no one reads.

For Open Source Maintainers

Cannot detect fake contributors, AI spam committers, or reputation farmers degrading community health. GitHub stars are easily manufactured.

For Freelancers

Reputation resets across platforms. 5 years on Fiverr earns zero credibility on Upwork. Every new platform means starting from scratch.

For Investors

Evaluating technical founders is guesswork without access to their actual output history. Pitch decks are polished; execution records are hidden.

For Grant Bodies & Accelerators

Applications are self-reported. No external verification of claimed credentials, past work, or community standing exists in a standardised form.

3. Product Architecture

A TrustGraph profile is not a portfolio. It is a living, evidence-backed record of capability that replaces claims with proof.

3.1 The Four Layers

1	Capability Layer	Shows verified skills — Frontend, Backend, AI, Cloud, Leadership, Research — all grounded in public evidence, not self-assessment.
2	Evidence Layer	Every claim links to verifiable proof. 240 merged PRs. 37 repositories. 4 public talks. 3 hackathon wins. The record is the profile.
3	Consistency Layer	Tracks long-term activity across years. Sustained effort weighs more than recent bursts. This is the anti-gaming mechanism.
4	Trust Layer	Verified vs. unverified claim ratio. Every score is explainable — positive and negative signals listed explicitly.

3.2 Trust Score Dimensions



Score Methodology (High Level)

The score is not a single number computed in one pass. Each dimension is scored independently, then combined with decay-adjusted weighting:

- Evidence Depth: counts and quality of verifiable artefacts (PRs, papers, talks, projects). Normalised by category so a conference keynote ≠ a lightning talk.
- Consistency: rolling 90-day and 365-day activity windows. A score earned over 3 years decays at 0.5% per month of inactivity beyond 90 days.
- Impact Signals: downstream effects — stars, forks, citations, project adoption, talk replay counts. Adjusted for platform inflation (GitHub stars are discounted vs. citations).
- Peer Verification: structured reference signals from verified co-workers, collaborators, or clients. Weighted by the verifier's own trust score.

Decay Rule

Old evidence is not deleted — it is decay-adjusted. A 2019 PR still counts, but at ~60% of the weight of a 2024 PR at the same quality level.

3.3 Score Sample

Trust Score: 88 / 100

- + Open source contributions (42 merged PRs, 3 maintained repos)
- + Speaking experience (2 conference talks, 1 workshop)
- + Multi-year consistency (active since 2021 with no gap > 60 days)
- 3 claimed projects without public repository or demo link
- Peer verification pending (0 of 3 references confirmed)

4. Onboarding & First-Run Experience

The first 5 minutes determine whether a user stays. A score of 0 on signup is a conversion killer. The onboarding flow must demonstrate value before the user has done any manual work.

4.1 The Bootstrapped Profile

On signup, TrustGraph immediately begins building a partial profile from any publicly available data tied to the user's email or connected username. Before the user answers a single question, we show them what we already know:

Shadow Profile Reveal	"We found 3 data sources linked to your GitHub handle. Your profile is 61% complete. Claim it below." — This is the hook. No empty state.
--------------------------	---

4.2 Onboarding Steps

Step	What Happens	Time
1 — Connect GitHub	OAuth connect. Immediately shows: repos, merged PRs, contribution graph, languages. Score appears in real-time as data loads.	< 30 sec
2 — Add one more source	Prompt: 'Add Stack Overflow, Devpost, or LinkedIn to increase your score.' Each connection adds 8–15 points visibly.	1–2 min
3 — Score reveal	First full trust score shown with dimension breakdown. The 'explainability panel' shows exactly what drove each dimension.	Instant
4 — Copy your Trust Passport	One-click copy of trustgraph.com/handle. Prompt: 'Share this wherever you'd normally share your resume.'	5 sec
5 — Invite for peer verification	Send 3 verification requests to colleagues. They confirm structured signals. Each confirmation adds ~5 points.	1 min

4.3 Empty State Handling

If a user has no detectable public data (rare for developers, possible for early career), we do not show a score of 0. We show a progress tracker:

- 'Connect GitHub to unlock Evidence score' — with a live preview of what their score could look like
- 'Answer 3 quick questions to bootstrap your profile while we gather evidence'

- Score appears as 'Estimated: 45–60' range until evidence confirms it — honesty over false precision

5. Full Feature Set

Twelve distinct features spanning personal profiles, team intelligence, enterprise tooling, and viral distribution surfaces.

5.1 Trust Passport

A portable reputation link — trustgraph.com/username — that users share in place of a resume link. The public view shows a curated trust summary without requiring the viewer to log in.

Unauthenticated View

Non-logged-in visitors see: Trust Score, top 3 verified capabilities, consistency timeline, and evidence count. Full dimension breakdown requires login. This is the viral teaser.

5.2 Embed Widget & Badges

A distributable trust badge that users embed on GitHub READMEs, personal websites, LinkedIn banners, email signatures, and portfolio sites.

- SVG badge: 'TrustGraph Verified — Score: 84' — one line of markdown
- Full widget: embeddable card with score, top capabilities, and evidence count
- Dynamic: badge updates automatically when score changes
- Every badge is a backlink and an impression — this is the primary viral surface

README Badge (Markdown)

 — drops into any GitHub README in 10 seconds.

5.3 Shadow Profiles

TrustGraph silently builds partial profiles for high-signal individuals — top GitHub contributors, paper authors, hackathon winners, conference speakers — before they sign up.

- When they arrive: 'We have 71% of your profile ready. Claim it in 30 seconds.'
- Removes cold-start friction entirely for the most valuable early users
- Legal basis: only public data; claim flow requires identity confirmation
- Seeding strategy: start with top 10,000 GitHub contributors by merged PR count

5.4 Trust Decay & Activity Alerts

The score is alive. It decays with inactivity and grows with output. This creates habitual re-engagement — the fitness tracker model applied to professional reputation.

- 'Your consistency score dropped 3 points — no public activity in 91 days.'
- 'You shipped 2 projects this month — your Evidence score increased by 6 points.'
- Weekly digest email: score delta, top movers, what to do next
- This is the retention engine — a reason to check back that isn't manufactured

5.5 Comparative Insights

Users understand not just their absolute score but their relative standing by dimension and category — the most actionable form of feedback.

You are in the top 11% of AI engineers by Evidence Depth

but your Consistency score is in the bottom 28% of engineers with 3+ years experience.

3 consistent months of public activity would move you into the top 40% overall.

5.6 Trust Timeline

A chronological view of when evidence milestones were reached. Turns a number into a story. Stories are shareable; numbers are not.

- First public commit → First merged PR → First hackathon win → First conference talk → Score crosses 80
- Shareable as a link or embeddable graphic
- The narrative arc creates pride and aspiration simultaneously

5.7 Anonymous Recruiter Signals

Pro tier feature. When a recruiter views a profile (even without login, inferred by traffic pattern), the profile owner sees:

- '4 people from fintech companies viewed your profile this week.'
- 'Your profile was included in a recruiter shortlist on Tuesday.'
- This mirrors LinkedIn's 'who viewed your profile' mechanic and drives Pro conversion immediately

5.8 Team Trust Score

Aggregated trust evaluation for entire teams. Useful for investors, accelerators, and hackathon organisers.

Example Output

Team: Acme AI | Engineering: 91 | Execution: 86 | Consistency: 88 | Peer Verification: 79

- Team score is not an average — it is a weighted composition that accounts for role coverage and complementarity
- A team of 3 engineers and 0 designers scores lower on 'coverage' than a balanced team
- Team score decays if key members leave (with a grace period and warning)

5.9 Startup Trust Pages

A startup profile showing founders, verified evidence, growth milestones, and public achievements — replacing pitch decks with proof-backed credentials for investor and accelerator audiences.

5.10 Community Trust

For GDGs, open source maintainers, DAOs, and hackathon organisers — understand contributor quality and detect farming behaviour before it degrades community health.

- Contributor quality score: depth-adjusted (not just commit count)
- Farming detection: velocity anomalies, low-signal commit patterns, new account spikes
- Community leaderboard: publicly visible, opt-in, drives aspirational pull

5.11 Reference Verification

Structured endorsement signals instead of long-form recommendations. Contacts confirm specific capabilities that feed directly into the Peer Verification dimension.

- Verifier confirms: skill area, collaboration context, time period, and confidence level
- Verifier's own trust score weights their endorsement — a high-trust verifier's signal counts more
- Verification requests expire after 14 days — keeps the graph fresh

5.12 Appeal & Dispute Flow

A legally necessary and trust-building feature for any product used in hiring contexts.

- 'Flag this evidence as incorrect / no longer mine' — on any evidence item
- 'Remove this data source' — with pre-removal score impact shown before confirmation
- Human review queue for enterprise clients with SLA commitment (48h standard, 4h enterprise)
- Outcome logged publicly: 'Evidence disputed — under review' keeps the graph honest

6. Privacy & Data Architecture

Privacy is not a footnote. The answers to these questions define what kind of product TrustGraph is.

6.1 Data Visibility Model

Scenario	Who Can See What
User's own profile	Full view — all evidence, all dimensions, all score history
Logged-in viewer	Public trust summary — score, top capabilities, consistency bar, evidence count. Detailed breakdown requires subject's consent.
Recruiter (paid tier)	Same as logged-in viewer unless user opts into 'recruiter-discoverable' mode. Not discoverable by default.
Unauthenticated visitor with link	Teaser view — score, top 3 capabilities, evidence count. No full breakdown. CTAs to sign up.
API caller (B2B)	Score band and binary threshold response only, unless user has explicitly opted into full API disclosure.

6.2 User Control Principles

- Every data source connection is explicit and reversible. Disconnecting a source shows score impact before confirmation.
- 'Private mode' — profile exists and score is computed, but the profile link returns 404 to non-permitted viewers. The Trust Passport still works when shared directly.
- Right to erasure (GDPR Art. 17): deleting account removes profile within 30 days. Cached shadow profile data purged within 72 hours of claim + delete.
- Users can opt out of shadow profile creation by emailing a designated address or via a public opt-out form.

6.3 What TrustGraph Never Does

- Never allows a recruiter to discover a user who has not opted into recruiter visibility
- Never sells raw profile data or evidence records to third parties
- Never allows 'boosted' or 'featured' placement that money can buy — scores are evidence-only
- Never builds scores on private repository data or non-public activity without explicit consent

**Principal-Agent
Commitment**

TrustGraph serves users first, recruiters second. Recruiters pay for access to a fair, evidence-based market — not for preferential treatment.

7. Technical Architecture (Product View)

Key technical decisions that define product behaviour, API capabilities, and enterprise readiness.

7.1 Data Freshness

Source Type	Update Cadence	Method
GitHub activity	Daily	GitHub API polling + webhook for owned repos
Devpost / hackathons	Weekly	Scheduled scrape + user-triggered refresh
Stack Overflow	Weekly	SO API polling
Conference / talks	Monthly or on-demand	Manual claim + URL verification
Peer verifications	Real-time	Event-driven on verification submission
Overall trust score	Recomputed daily	Async background job; user sees last computed + 'updating' state

7.2 API Design

- Base endpoint: GET /v1/trust-score?handle={username}
- Response modes: full (opted-in users), band (75–85, requires lower consent), binary (above/below threshold, minimum consent)
- Webhook: POST to registered URL when score changes by >5 points
- SDK: JavaScript and Python SDKs in Phase 1; Go and Ruby in Phase 2
- SLA: 99.9% uptime for API tier. Response < 200ms for cached scores.

7.3 Anti-Gaming Measures

- Multi-dimensional scoring: optimising one dimension (e.g., commit count) without others (impact, consistency) does not move the overall score meaningfully
- Velocity anomaly detection: sudden spike in any evidence type triggers a review hold on that dimension
- New account signal: evidence from accounts < 90 days old is quarantined with reduced weighting
- AI contribution detection (Phase 1 research item): commit pattern analysis to flag low-signal AI-generated contributions
- Cross-source corroboration: a GitHub project with no external mentions, stars, forks, or collaborators gets a low impact score regardless of commit count

8. Go-To-Market & Growth Strategy

Phase 1 success depends entirely on getting the right first 1,000 profiles. Quality of early users defines the benchmark that all future users are measured against.

8.1 Seeding Strategy — Phase 1

- Target: top 2,000 GitHub contributors in AI, web, and systems categories by merged PR count
- Method: cold outreach via GitHub profile email or Twitter/X DM — 'We built your profile. Want to claim it?'
- Partner: reach out to 5–10 popular developer newsletters (TLDR, Pointer, Changelog) for a launch feature
- Hackathon activation: sponsor one major hackathon (MLH, Devfolio) in Month 1. All winners get TrustGraph profiles created automatically.
- Community seeding: post in r/programming, Hacker News Show HN, and dev.to on launch day with a real profile demo

8.2 Viral Mechanics (Designed, Not Hoped For)

Surface	Trigger	Conversion Goal
Unauthenticated profile view	User shares Trust Passport link	Viewer clicks 'Verify my profile too'
GitHub README badge	User adds badge to README	Repo visitor clicks badge → landing page
Email signature badge	User adds score badge to email sig	Email recipient clicks → profile view
Hackathon submission	User links TrustGraph in project	Judges / other participants view and sign up
Recruiter discovery	Recruiter searches, finds profile	Subject sees 'someone viewed you' → shares link

8.3 Demand-Side Seeding (Recruiter)

The platform needs demand (recruiters, investors) before supply (profiles) saturates. Chicken-and-egg mitigation:

- Free recruiter accounts in Phase 1 — let them search existing profiles at no cost. Capture intent and feedback.
- 'Trust-verified shortlists' — a free tool recruiters can use to verify a candidate they already found elsewhere. Drives recruiter adoption without needing full search.

- Partner with 2–3 hiring tools (Ashby, Greenhouse, Lever) for early API integration. Their users bring demand.

9. Business Model

Tier	Price	Target	Core Value
Free	No cost	All professionals	Trust profile, basic score, limited analytics, Trust Passport link
Pro	\$15/month	Active professionals	Advanced insights, score history, comparisons, recruiter signal alerts, full timeline
Recruiter	\$149–499/month	Talent teams	Search by trust score, skill, location, activity. Shortlist exports. Candidate verification.
Startup	\$199/month	Founding teams	Team Trust Score, Startup Trust Pages, investor-ready reports, team composition analysis
API — Band	\$0.05/call	HR tech, ATS platforms	Score band response (e.g., 75–85). No raw score. Minimum consent model.
API — Full	\$0.20/call	Enterprise, regulated hiring	Full score + dimension breakdown. Requires user opt-in. SLA guaranteed.
White Label	Custom / \$2k+/mo	Universities, accelerators	Branded trust scoring system. Custom data sources. Dedicated support.

Biggest Long-Term Revenue	API revenue. Every ATS, hiring platform, freelance marketplace, and grant portal embeds one API call. Network effect multiplies this.
---------------------------	---

Revenue Projections (Conservative)

Phase	Timeframe	Primary Revenue	MRR Target
1 — Developer Foundation	M1–M6	Pro subscriptions + early Recruiter plans	\$8,000–15,000
2 — Designer Expansion	M7–M12	Pro + Recruiter (design studios)	\$25,000–45,000
3 — Researcher Network	M13–M18	White label (universities) + API beta	\$60,000–90,000
4 — Founder & Startup Layer	M19–M30	Startup plan + investor dashboard + API growth	\$150,000–250,000
5 — General Professional	M31+	API at scale + enterprise white label	\$500,000+

10. Risks & Mitigations

Risk	Impact	Mitigation
Score gaming / Goodhart's Law	High	Multi-dimensional scoring; no single axis can move the total score meaningfully. Velocity anomaly detection freezes suspect dimensions.
Privacy & data sovereignty (GDPR, CCPA)	High	Explicit consent at every source connection. Right to erasure within 30 days. Shadow profile opt-out mechanism. Data residency options for enterprise.
Bias in scoring models	High	Full explainability (no black box). Regular third-party bias audits. Appeal mechanism. No hidden weighting adjustments.
Platform API access revoked (GitHub, SO)	Medium	Multi-source redundancy — score continues with remaining sources. Historical data cached. 30-day notice fallback period.
AI-generated fake contributions	High	Commit pattern analysis. Cross-source corroboration requirement. Quarantine + review for velocity anomalies.
Legal exposure in hiring use cases	High	API returns score bands, not raw numbers. No protected-class data in scoring. Legal review before enterprise HR integrations.
Cold-start for non-developer segments	Medium	Phase expansion is evidence-density-gated. Segment only entered when 3+ verifiable public data sources exist.
LinkedIn blocking data access	Medium	LinkedIn is secondary; primary sources are direct platform APIs. Score functional without LinkedIn.
Recruiter-user principal-agent conflict	High	Structural: user opt-in required for recruiter discovery. No paid boosting. No 'featured' scores. Policy published and auditable.

11. Competitive Landscape

Platform	What It Does	Core Limitation	TrustGraph Advantage
LinkedIn	Professional profile + network	100% self-reported. Endorsements are meaningless. Optimised for recruiters, not truth.	Evidence-backed. Not self-reported. Portable outside LinkedIn.
GitHub	Code repository + contribution graph	Code only. No cross-platform signal. No trust score. No portability.	Cross-platform. Scores non-code work. Portable link.
read.cv	Minimal professional profile	Still self-reported. No verification. No scoring. No API.	Verified. Scored. API-first. Evidence-driven.
Contra	Freelancer independent work platform	Platform-locked. No portability. Creator economy focus only.	Platform-agnostic. Every profession. Portable.
Devpost	Hackathon submissions & profiles	Hackathon-only signal. No career-wide view. No score.	Aggregates Devpost as one signal among many.
Google Scholar	Research citation tracking	Research only. No other professional signals.	Aggregates Scholar as one signal among many.
HireQuotient / Workera	Skill assessment platforms	Test-based, not evidence-based. High friction. No ongoing record.	Passive evidence gathering. No tests required.

Category Gap

No platform aggregates cross-discipline professional reputation into a portable, evidence-based, explainable score. This category has no dominant winner.

12. Design Language & Direction

TrustGraph should feel like a trusted institution that was built natively for the internet — not a startup, not a social network, not a SaaS dashboard. Precise. Understated. Evidence-first.

12.1 Tone

- Precision over personality — let the evidence speak
- Confident, not boastful — a single large score number, no confetti
- Human warmth in the gaps — score explanation copy feels like a thoughtful colleague, not a machine
- Never gamified — no trophies, no streaks, no XP bars. This is a professional tool.

12.2 Reference Sites to Study

Site	Why to Study It
linear.app	Gold standard for developer-facing dark-mode product design. Precision, speed, minimal chrome. TrustGraph should feel this clean.
read.cv	Closest existing product conceptually. Study their profile layout, whitespace philosophy, and 'quiet confidence' tone.
cal.com	Open, transparent product philosophy translated into visual design. How a trust-forward brand looks.
clerk.com	Identity-adjacent infrastructure with great visual design. Relevant because TrustGraph is also trust infrastructure.
workos.com	B2B developer tool that feels trustworthy without feeling boring. Strong type hierarchy, clean tables.
contra.com	Freelancer identity platform. Study profile cards and how they handle the 'showcase without bragging' problem.
arc.dev	Developer hiring platform. Study developer profile design and skill verification UX.
stripe.com (docs)	The benchmark for trust-through-design. Study how Stripe communicates reliability visually — the underlying confidence, not the flashiness.

12.3 Typography

Role	Recommendation	Reason
Body & UI	Inter or Geist	Clean, technical, screen-optimised. Industry standard for dev tools.
Display / Score	Plus Jakarta Sans or Sora	Slightly geometric — gives the score number visual authority without being cold.

Evidence snippets	JetBrains Mono	Makes evidence items (PR counts, commit numbers) feel like log entries, not marketing copy.
-------------------	----------------	---

12.4 Colour System

- Background: #F8FAFC (near-white, not pure white — prevents eye fatigue, signals quality)
- Surface: #FFFFFF for cards
- Primary: #2563EB (confident blue — not teal, not indigo) — used only for primary actions and verified signals
- Verified: #059669 (muted green) — used ONLY for verified evidence, never decoration
- Text primary: #0F172A (near-black). Text secondary: #64748B (slate)
- Avoid: gradients everywhere, rainbow skill bars, neon, excessive animation

12.5 Score Visualisation

- The trust score is one large number — not a dial, not a pie chart, not a progress ring
- Dimension breakdown: thin horizontal bars below the score, labelled, using the verified green and slate palette only
- Delta: '+3 this month' in small type next to the score — the change is as important as the number
- No percentage bars for skills — they look made-up. Use evidence counts instead: '14 merged PRs in TypeScript' > '90% JavaScript'

12.6 What to Avoid

- Skill bars showing percentage (e.g. '90% JavaScript') — meaningless and unverified
- Trophy and badge overload — this is a professional tool, not a video game
- Dark patterns around score improvement — no 'complete your profile for 10 points' spam
- Hero images and cover photos — not Instagram. Whitespace is confidence, not emptiness
- Animations on everything — one or two purposeful micro-animations max (score counter on load, evidence item reveal)

13. Phase-Wise Execution Plan

Five phases across 36+ months. Each phase gates on evidence density — we only expand to a new segment when we have enough verifiable data sources to build meaningful scores.

PHASE 1	Developer Foundation <i>Timeframe: Months 1–6</i>	AUDIENCE Developers
-------------------	---	-------------------------------

Why First

Developers have the most verifiable public data of any professional category. GitHub, Stack Overflow, Devpost, and conference records are fully public and crawlable. Evidence collection cost is near-zero.

Phase 1 Data Sources

- GitHub — repositories, merged PRs, stars, forks, contributor graphs, language breakdown
- Devpost — hackathon wins and submissions, project records, team history
- Stack Overflow — accepted answers, reputation score, top answer tags
- Conference records — Sessionize, CFP.dev, YouTube talk archives
- Google Scholar / arXiv — research papers and citation counts (for developer-researchers)

Phase 1 Deliverables

- Core profile: Trust Passport link, capability layer, evidence layer, consistency timeline
- Trust Score MVP: four-dimension model with explainability panel
- Shadow profile system: top 10,000 GitHub contributors pre-built at launch
- Embed widget: GitHub README badge + personal site widget
- Viral loop: unauthenticated profile view with 'claim your profile' CTA
- Onboarding flow: OAuth connect → score reveal → passport copy (under 3 minutes)
- Free and Pro tiers live
- Trust decay and activity alerts (weekly digest email)
- Basic API: score band endpoint (private beta for 3–5 design partners)

Phase 1 Success Metrics

Metric	Target
--------	--------

Profiles created	10,000+
GitHub integrations active	> 80% of profiles
Viral coefficient (profile views → signups)	> 0.3
D30 retention	> 40%
Pro tier conversion	> 5%
README badge embeds	> 500
API design partner integrations	3–5

PHASE
2

Designer Expansion
Timeframe: Months 7–12

AUDIENCE
Designers

Why Second

Designers have substantial verifiable public output — Dribbble, Behance, Figma Community, and design conference records are public and indexable. Designer hiring is broken in the same way developer hiring is. And cross-discipline team scores (dev + design) become more valuable from this phase onward.

Phase 2 New Data Sources

- Dribbble — follower-adjusted engagement, shot quality, case study depth indicators
- Behance — project views, appreciations, Adobe Certified credentials
- Figma Community — template downloads, plugin installs, remix counts
- Design conference talks and mentorship records
- UX case study verification: shipped products vs. concepts (App Store / web presence cross-check)

Phase 2 Deliverables

- Designer-specific capability dimensions: Craft, UX Depth, Shipping Record, Visual Communication
- Portfolio evidence layer: visual work thumbnails linking to live sources
- Cross-role team trust: dev + design team scores now available
- Recruiter Plan live: with Designer filters (craft score, portfolio depth, shipped products)
- Community trust beta: for design communities and open source design systems

Phase 2 Success Metrics

- Designer profiles: 30%+ of new signups in M9–M12

- First Recruiter Plan subscriptions from design studios or product teams
- Cross-discipline team scores: 100+ team profiles created
- MRR target: \$25,000–45,000

PHASE 3	Researcher Network <i>Timeframe: Months 13–18</i>	AUDIENCE Researchers
------------------------------	--	---

Why Third

Researchers have the richest existing verification infrastructure of any profession — citation counts, peer review records, conference programme committee roles, and institutional affiliations. TrustGraph supplements these with cross-platform signals.

Phase 3 New Data Sources

- Google Scholar — h-index, citations, co-author networks, publication venues
- Semantic Scholar / OpenAlex — paper impact, field influence, citation velocity
- arXiv / bioRxiv / SSRN — preprint velocity, community engagement, download counts
- Conference programme committee memberships and review invitations
- Grant records — NSF, NIH, Wellcome, ERC public award databases
- Patent records — USPTO, EPO public filing data

Phase 3 Deliverables

- Research trust score: citation-adjusted impact, venue prestige weighting, h-index normalised by field
- Institution verification layer: affiliated email domain cross-check
- Grant application trust reports: TrustGraph profile as a credibility supplement to grant applications
- University white-label offering: branded trust scoring for student and faculty profiles
- API general availability: open to all developers, public rate limits, full documentation

Phase 3 Success Metrics

- 5+ university partners on white-label plan
- Research profiles with Google Scholar integration: 20,000+
- First API revenue from grant bodies or accelerators
- MRR target: \$60,000–90,000

Founder & Startup Layer

Timeframe: Months 19–30

Why Fourth

By Phase 4 the trust graph is rich enough to evaluate founding teams holistically. Founder trust draws from developer, designer, and researcher scores already established — plus execution signals specific to company building. The product now serves both sides of the funding relationship.

Phase 4 New Data Sources

- Crunchbase — funding rounds, co-investor records, company growth signals
- Product Hunt — launch reception, upvotes, community comment sentiment
- AngelList — deal history, investor reference signals
- GitHub org activity — team velocity, shipping cadence, contributor retention
- Reference verification from co-founders, early employees, first customers
- Accelerator / incubator records (with partner data-sharing agreements)

Phase 4 Deliverables

- Startup Trust Pages: investor-ready, proof-backed company profiles
- Team Trust Score for entire founding teams with role coverage analysis
- Investor dashboard: filter startups by team trust threshold, track portfolio company scores
- Startup Plan tier: team verification, investor-ready trust reports, team composition gap analysis
- Accelerator partnerships: YC, Techstars, and equivalent program integrations
- Cohort trust scoring: accelerators use TrustGraph to score applicant teams in batch

Phase 4 Success Metrics

- Startup Trust Pages: 1,000+ company profiles
- Accelerators using TrustGraph for cohort selection: 3+
- Investor dashboard MRR: \$50,000+
- Total MRR target: \$150,000–250,000

General Professional Network

Timeframe: Month 31+

The Horizontal Expansion

Phase 5 opens TrustGraph to all professional categories — finance, law, medicine, marketing, operations — using the trust infrastructure, scoring models, and API layer built across the previous four phases. The product transitions from a developer tool to professional infrastructure.

Phase 5 New Data Sources

- Professional licensing records — bar association, medical licensing, CPA public registries
- LinkedIn professional history: tenure signals, endorsements (weighted by endorser trust score)
- Upwork / Fiverr / Toptal: cross-platform freelancer rating aggregation
- Salesforce / HubSpot verified revenue records (opt-in for sales professionals)
- Community leadership: published articles, event organisation, certification bodies
- Government contractor records (public procurement databases)

Phase 5 Deliverables

- Universal trust score: profession-agnostic, horizontally applicable
- API v2: enterprise integrations for ATS, HRMS, and ERP platforms
- White-label for staffing agencies, enterprises, and government bodies
- TrustGraph resume replacement: official positioning shift in marketing
- Mobile app: iOS and Android profile management and alert system
- Global data residency: EU, US, Asia-Pacific infrastructure for enterprise compliance

14. Phase-Wise Build Timeline

Sprint-level timeline for the first 18 months (Phases 1–3). Each sprint is 2 weeks. Owner column shows the function responsible — adapt to your actual team structure.

Phase 1 — Developer Foundation (Months 1–6)

Sprint	Deliverable / Task	Owner
W1–W2	Project scaffolding: monorepo setup, CI/CD pipeline, auth system (OAuth + magic link), database schema v1	Eng
W3–W4	GitHub OAuth integration: connect account, ingest repos, PRs, contributor graph. Score engine v1: evidence dimension only	Eng
W5–W6	Profile page v1: Trust Passport link, capability display, evidence list, consistency timeline bar	Eng + Design
W7–W8	Trust Score full model: all 4 dimensions live. Explainability panel. Score delta ('this month')	Eng
W9–W10	Shadow profile system: ingest top 10,000 GitHub contributors. Claim flow (email / GitHub handle verification)	Eng
W11–W12	Onboarding flow: connect → score reveal → passport copy → peer verification invite. Under 3-minute target	Eng + Design
W13–W14	Embed widget v1: SVG badge for GitHub README. Dynamic score update on re-render	Eng
W15–W16	Stack Overflow + Devpost integrations. Score refresh pipeline (daily GitHub, weekly SO/Devpost)	Eng
W17–W18	Pro tier: Stripe billing, advanced analytics, score history graph, comparative insights (percentile view)	Eng + Product
W19–W20	Activity alerts: weekly digest email, score decay notifications, evidence milestone celebrations	Eng + Marketing
W21–W22	Trust Decay engine: decay coefficients live. Inactivity alerts. Score recalculation on reconnect.	Eng
W23–W24	API private beta: score band endpoint, API key management, 3–5 design partner integrations. Launch prep.	Eng + Product

Phase 2 — Designer Expansion (Months 7–12)

Sprint	Deliverable / Task	Owner
--------	--------------------	-------

W25–W26	Dribbble + Behance OAuth integrations. Designer capability dimension model (Craft, UX Depth, Shipping Record)	Eng
W27–W28	Figma Community integration. Portfolio evidence layer: visual work thumbnails with source links	Eng + Design
W29–W30	Cross-role team trust: combined dev + design team score. Team profile page v1	Eng + Product
W31–W32	Recruiter Plan: Stripe billing, talent search UI, trust score + skill + location filters, shortlist exports	Eng + Product
W33–W34	Community trust beta: maintainer dashboard, contributor quality scores, farming detection v1	Eng
W35–W36	Appeal & dispute flow: flag evidence, remove source with impact preview, human review queue	Eng + Ops
W37–W38	Anonymous recruiter signal feature (Pro): 'X people from Y industry viewed your profile this week'	Eng
W39–W40	Trust Timeline feature: milestone history, shareable graphic, embeddable card	Eng + Design
W41–W42	API GA (general availability): public rate limits, full docs site, JavaScript + Python SDKs, webhook support	Eng + DevRel
W43–W44	Mobile-responsive profile redesign. Performance audit. Core Web Vitals pass.	Eng + Design
W45–W46	Velocity anomaly detection v1: flag sudden spikes in evidence count for review. Anti-gaming hardening.	Eng
W47–W48	Phase 3 prep: Google Scholar API integration scaffolding. University partner conversations begin.	Eng + BD

Phase 3 — Researcher Network (Months 13–18)

Sprint	Deliverable / Task	Owner
W49–W50	Google Scholar integration live. Research trust dimension: citation-adjusted impact, h-index, venue prestige weighting	Eng
W51–W52	Semantic Scholar + OpenAlex integration. arXiv / bioRxiv preprint velocity signals	Eng
W53–W54	Institution verification layer: affiliated email domain cross-check. Institutional badge on profiles	Eng + Ops
W55–W56	Grant application trust report: PDF export of researcher trust profile formatted for grant submission	Eng + Design
W57–W58	White-label v1: university branding, custom domain, isolated data environment, admin dashboard	Eng + Product
W59–W60	University partner onboarding: 5-partner target. Dedicated onboarding flow and support SLA	BD + Ops

W61–W62	API billing v2: per-call metering, dashboard, overage alerts, enterprise contract support	Eng + Finance
W63–W64	Go SDK + Ruby SDK. API integration guides for Greenhouse, Lever, Ashby (ATS platforms)	Eng + DevRel
W65–W66	Bias audit v1: third-party review of scoring model across gender, institution type, geography	Product + External
W67–W68	GDPR compliance hardening: right to erasure automation, data processing agreements, DPA templates for enterprise	Eng + Legal
W69–W70	Phase 4 prep: Crunchbase + Product Hunt integration scaffolding. Investor dashboard wireframes.	Eng + Product
W71–W72	Phase 3 retrospective, score model recalibration, Phase 4 kickoff.	All

Phases 4–5 — High-Level Milestone Overview

Detailed sprint plans for Phases 4 and 5 are produced at the end of Phase 3, informed by what we learned. Key milestones:

M19

Startup layer launch

- Crunchbase + Product Hunt integrations live
- Startup Trust Pages GA
- Team Trust Score for founding teams
- Investor dashboard private beta

M22

Accelerator partnerships

- First 3 accelerators using TrustGraph for cohort scoring
- Cohort trust scoring bulk API endpoint
- Startup Plan billing live

M28

Enterprise readiness

- SOC 2 Type II audit initiated
- SSO / SAML for enterprise white-label
- Dedicated CSM tier for enterprise accounts

M31

General professional expansion

- Universal score model (non-developer dimensions)
- API v2 with ATS/HRMS integrations
- Mobile app launch (iOS + Android)

M36

Infrastructure positioning

- TrustGraph officially positioned as reputation infrastructure
- Government and regulated sector white-label pilots
- Resume replacement campaign launched

15. Defensibility & Network Effects

The Moat

Not AI — AI is replaceable. The moat is the trust graph, historical evidence accumulation, and the verification network. Similar to why LinkedIn is hard to displace.

Flywheel

**More users → More endorsements → More evidence → Better trust models →
More recruiters → More users**

Long-Term Vision

10 years from now, professionals include TrustGraph alongside a resume and LinkedIn. Then eventually TrustGraph replaces the resume.

The end state is not a profile website. TrustGraph becomes the global reputation infrastructure layer sitting underneath hiring, freelancing, startup formation, grants, communities, and open source collaboration. A category with no dominant winner yet.

The end state is not a profile website.

TrustGraph becomes the global reputation infrastructure layer.